



Red Hat Ansible Automation Platform 2.5

Installing on OpenShift Container Platform

Install and configure Ansible Automation Platform operator on OpenShift Container Platform

Red Hat Ansible Automation Platform 2.5 Installing on OpenShift Container Platform

Install and configure Ansible Automation Platform operator on OpenShift Container Platform

Legal Notice

Copyright © 2024 Red Hat, Inc.

The text of and illustrations in this document are licensed by Red Hat under a Creative Commons Attribution–Share Alike 3.0 Unported license ("CC-BY-SA"). An explanation of CC-BY-SA is available at

<http://creativecommons.org/licenses/by-sa/3.0/>

. In accordance with CC-BY-SA, if you distribute this document or an adaptation of it, you must provide the URL for the original version.

Red Hat, as the licensor of this document, waives the right to enforce, and agrees not to assert, Section 4d of CC-BY-SA to the fullest extent permitted by applicable law.

Red Hat, Red Hat Enterprise Linux, the Shadowman logo, the Red Hat logo, JBoss, OpenShift, Fedora, the Infinity logo, and RHCE are trademarks of Red Hat, Inc., registered in the United States and other countries.

Linux[®] is the registered trademark of Linus Torvalds in the United States and other countries.

Java[®] is a registered trademark of Oracle and/or its affiliates.

XFS[®] is a trademark of Silicon Graphics International Corp. or its subsidiaries in the United States and/or other countries.

MySQL[®] is a registered trademark of MySQL AB in the United States, the European Union and other countries.

Node.js[®] is an official trademark of Joyent. Red Hat is not formally related to or endorsed by the official Joyent Node.js open source or commercial project.

The OpenStack[®] Word Mark and OpenStack logo are either registered trademarks/service marks or trademarks/service marks of the OpenStack Foundation, in the United States and other countries and are used with the OpenStack Foundation's permission. We are not affiliated with, endorsed or sponsored by the OpenStack Foundation, or the OpenStack community.

All other trademarks are the property of their respective owners.

Abstract

This guide provides procedures and reference information for the supported installation scenarios for the Red Hat Ansible Automation Platform operator on OpenShift Container Platform.

Table of Contents

PREFACE	4
PROVIDING FEEDBACK ON RED HAT DOCUMENTATION	5
CHAPTER 1. PLANNING YOUR RED HAT ANSIBLE AUTOMATION PLATFORM OPERATOR ON RED HAT OPENSIFT CONTAINER PLATFORM	6
1.1. ABOUT ANSIBLE AUTOMATION PLATFORM OPERATOR	6
1.2. OPENSIFT CONTAINER PLATFORM VERSION COMPATIBILITY	6
1.3. SUPPORTED INSTALLATION SCENARIOS FOR RED HAT OPENSIFT CONTAINER PLATFORM	6
1.4. CUSTOM RESOURCES	7
1.5. ADDITIONAL RESOURCES	7
CHAPTER 2. INSTALLING THE RED HAT ANSIBLE AUTOMATION PLATFORM ON RED HAT OPENSIFT CONTAINER PLATFORM	8
CHAPTER 3. INSTALLING RED HAT ANSIBLE AUTOMATION PLATFORM OPERATOR FROM THE OPENSIFT CONTAINER PLATFORM CLI	9
3.1. PREREQUISITES	9
3.2. SUBSCRIBING A NAMESPACE TO AN OPERATOR USING THE OPENSIFT CONTAINER PLATFORM CLI	9
3.3. FETCHING PLATFORM GATEWAY LOGIN DETAILS FROM THE OPENSIFT CONTAINER PLATFORM CLI	11
3.3.1. Fetching the platform gateway web address	11
3.3.2. Fetching the platform gateway password	11
3.4. ADDITIONAL RESOURCES	12
CHAPTER 4. CONFIGURING THE RED HAT ANSIBLE AUTOMATION PLATFORM OPERATOR ON RED HAT OPENSIFT CONTAINER PLATFORM	13
4.1. LINKING YOUR COMPONENTS TO THE PLATFORM GATEWAY	13
4.2. ACCESSING THE PLATFORM GATEWAY	14
4.3. DEPLOYING THE PLATFORM GATEWAY WITH EXISTING ANSIBLE AUTOMATION PLATFORM COMPONENTS	15
4.4. FREQUENTLY ASKED QUESTIONS ON PLATFORM GATEWAY	16
CHAPTER 5. CONFIGURING AUTOMATION CONTROLLER ON RED HAT OPENSIFT CONTAINER PLATFORM WEB CONSOLE	17
5.1. PREREQUISITES	17
5.1.1. Configuring your controller image pull policy	17
5.1.2. Configuring your controller LDAP security	18
5.1.3. Configuring your automation controller operator route options	19
5.1.4. Configuring the ingress type for your automation controller operator	20
5.2. CONFIGURING AN EXTERNAL DATABASE FOR AUTOMATION CONTROLLER ON RED HAT ANSIBLE AUTOMATION PLATFORM OPERATOR	21
5.3. FINDING AND DELETING PVCS	23
5.4. ADDITIONAL RESOURCES	23
CHAPTER 6. CONFIGURING AUTOMATION HUB ON RED HAT OPENSIFT CONTAINER PLATFORM WEB CONSOLE	25
6.1. PREREQUISITES	25
6.1.1. Storage options for Ansible Automation Platform Operator installation on Red Hat OpenShift Container Platform	25
6.1.1.1. Provisioning OCP storage with ReadWriteMany access mode	25
6.1.1.2. Configuring object storage on Amazon S3	26
6.1.1.3. Configuring object storage on Azure Blob	26
6.1.2. Configure your automation hub operator route options	27

6.1.3. Configuring the ingress type for your automation hub operator	28
6.2. CONFIGURING LDAP AUTHENTICATION FOR ANSIBLE AUTOMATION HUB ON OPENSIFT CONTAINER PLATFORM	29
6.3. FINDING THE AUTOMATION HUB ROUTE	30
6.4. CONFIGURING AN EXTERNAL DATABASE FOR AUTOMATION HUB ON RED HAT ANSIBLE AUTOMATION PLATFORM OPERATOR	30
6.4.1. Enabling the hstore extension for the automation hub PostgreSQL database	32
6.5. FINDING AND DELETING PVCS	33
6.6. ADDITIONAL CONFIGURATIONS	33
6.7. ADDITIONAL RESOURCES	33
CHAPTER 7. DEPLOYING CLUSTERED REDIS ON RED HAT ANSIBLE AUTOMATION PLATFORM OPERATOR	35
CHAPTER 8. USING RED HAT SINGLE SIGN-ON OPERATOR WITH AUTOMATION HUB	36
8.1. CREATING A KEYCLOAK INSTANCE	36
8.2. CREATING A KEYCLOAK REALM FOR ANSIBLE AUTOMATION PLATFORM	37
8.3. CREATING A KEYCLOAK CLIENT	38
8.4. CREATING A KEYCLOAK USER	41
8.5. INSTALLING THE ANSIBLE AUTOMATION PLATFORM OPERATOR	42
8.6. CREATING A RED HAT SINGLE SIGN-ON CONNECTION SECRET	42
8.7. INSTALLING AUTOMATION HUB USING THE ANSIBLE AUTOMATION PLATFORM OPERATOR	43
8.8. DETERMINING THE AUTOMATION HUB ROUTE	44
8.9. UPDATING THE RED HAT SINGLE SIGN-ON CLIENT	45
8.10. ADDITIONAL RESOURCES	46
CHAPTER 9. ADDING EXECUTION NODES TO RED HAT ANSIBLE AUTOMATION PLATFORM OPERATOR ..	47
CHAPTER 10. ANSIBLE AUTOMATION PLATFORM RESOURCE OPERATOR	49
10.1. RESOURCE OPERATOR OVERVIEW	49
10.2. USING RESOURCE OPERATOR	49
10.3. CONNECTING RESOURCE OPERATOR TO PLATFORM GATEWAY	49
10.4. CREATING A AUTOMATION CONTROLLER CONNECTION SECRET FOR RESOURCE OPERATOR	50
10.5. CREATING AN ANSIBLEJOB	50
10.6. CREATING A JOBTEMPLATE	51

PREFACE

Thank you for your interest in Red Hat Ansible Automation Platform. Ansible Automation Platform is a commercial offering that helps teams manage complex multi-tier deployments by adding control, knowledge, and delegation to Ansible-powered environments.

This guide helps you to understand the installation, migration and upgrade requirements for deploying the Ansible Automation Platform Operator on OpenShift Container Platform.

PROVIDING FEEDBACK ON RED HAT DOCUMENTATION

If you have a suggestion to improve this documentation, or find an error, you can contact technical support at <https://access.redhat.com> to open a request.

CHAPTER 1. PLANNING YOUR RED HAT ANSIBLE AUTOMATION PLATFORM OPERATOR ON RED HAT OPENSIFT CONTAINER PLATFORM

Red Hat Ansible Automation Platform is supported on both Red Hat Enterprise Linux and Red Hat OpenShift.

OpenShift operators help install and automate day-2 operations of complex, distributed software on Red Hat OpenShift Container Platform. The Ansible Automation Platform Operator enables you to deploy and manage Ansible Automation Platform components on Red Hat OpenShift Container Platform.

You can use this section to help plan your Red Hat Ansible Automation Platform installation on your Red Hat OpenShift Container Platform environment. Before installing, review the supported installation scenarios to determine which meets your requirements.

1.1. ABOUT ANSIBLE AUTOMATION PLATFORM OPERATOR

The Ansible Automation Platform Operator provides cloud-native, push-button deployment of new Ansible Automation Platform instances in your OpenShift environment. The Ansible Automation Platform Operator includes resource types to deploy and manage instances of automation controller and private automation hub. It also includes automation controller job resources for defining and launching jobs inside your automation controller deployments.

Deploying Ansible Automation Platform instances with a Kubernetes native operator offers several advantages over launching instances from a playbook deployed on Red Hat OpenShift Container Platform, including upgrades and full lifecycle support for your Red Hat Ansible Automation Platform deployments.

You can install the Ansible Automation Platform Operator from the Red Hat Operators catalog in OperatorHub.

For information about the Ansible Automation Platform Operator infrastructure topology see [Container topologies](#) in *Tested deployment models*.

1.2. OPENSIFT CONTAINER PLATFORM VERSION COMPATIBILITY

The Ansible Automation Platform Operator to install Ansible Automation Platform 2.5 is available on OpenShift Container Platform 4.9 and later versions.

Additional resources

- See the [Red Hat Ansible Automation Platform Life Cycle](#) for the most current compatibility details.

1.3. SUPPORTED INSTALLATION SCENARIOS FOR RED HAT OPENSIFT CONTAINER PLATFORM

You can use the OperatorHub on the Red Hat OpenShift Container Platform web console to install Ansible Automation Platform Operator.

Alternatively, you can install Ansible Automation Platform Operator from the OpenShift Container Platform command-line interface (CLI), **oc**. See [Installing Red Hat Ansible Automation Platform Operator from the OpenShift Container Platform CLI](#) for help with this.

After you have installed Ansible Automation Platform Operator you must create an **Ansible Automation Platform** custom resource (CR). This enables you to manage Ansible Automation Platform components from a single unified interface known as the platform gateway. As of version 2.5, you must create an Ansible Automation Platform CR, even if you have an existing automation controller, automation hub, or Event-Driven Ansible, components.

If existing components have already been deployed, you must specify these components on the Ansible Automation Platform CR. You must create the custom resource in the same namespace as the existing components.

Supported scenarios	Supported scenarios with existing components
<ul style="list-style-type: none"> ● Ansible Automation Platform CR for blank slate install with automation controller, automation hub, and Event-Driven Ansible enabled ● Ansible Automation Platform CR with just automation controller enabled ● Ansible Automation Platform CR with just automation controller, automation hub enabled ● Ansible Automation Platform CR with just automation controller, Event-Driven Ansible enabled 	<ul style="list-style-type: none"> ● Ansible Automation Platform CR created in the same namespace as an existing automation controller CR with the automation controller name specified on the Ansible Automation Platform CR spec ● Same with automation controller and automation hub ● Same with automation controller, automation hub, and Event-Driven Ansible ● Same with automation controller and Event-Driven Ansible

1.4. CUSTOM RESOURCES

You can define custom resources for each primary installation workflows.

1.5. ADDITIONAL RESOURCES

- See [Understanding OperatorHub](#) to learn more about OpenShift Container Platform OperatorHub.

CHAPTER 2. INSTALLING THE RED HAT ANSIBLE AUTOMATION PLATFORM ON RED HAT OPENSIFT CONTAINER PLATFORM

Prerequisites

- You have installed the Red Hat Ansible Automation Platform catalog in OperatorHub.
- You have created a **StorageClass** object for your platform and a persistent volume claim (PVC) with **ReadWriteMany** access mode. See [Dynamic provisioning](#) for details.
- To run Red Hat OpenShift Container Platform clusters on Amazon Web Services (AWS) with **ReadWriteMany** access mode, you must add NFS or other storage.
 - For information about the AWS Elastic Block Store (EBS) or to use the **aws-ebs** storage class, see [Persistent storage using AWS Elastic Block Store](#).
 - To use multi-attach **ReadWriteMany** access mode for AWS EBS, see [Attaching a volume to multiple instances with Amazon EBS Multi-Attach](#).

Procedure

1. Log in to Red Hat OpenShift Container Platform.
2. Navigate to **Operators** → **OperatorHub**.
3. Search for Ansible Automation Platform and click **Install**.
4. Select an **Update Channel**:
 - **stable-2.x**: installs a namespace-scoped operator, which limits deployments of automation hub and automation controller instances to the namespace the operator is installed in. This is suitable for most cases. The stable-2.x channel does not require administrator privileges and utilizes fewer resources because it only monitors a single namespace.
 - **stable-2.x-cluster-scoped**: deploys automation hub and automation controller across multiple namespaces in the cluster and requires administrator privileges for all namespaces in the cluster.
5. Select **Installation Mode**, **Installed Namespace**, and **Approval Strategy**.
6. Click **Install**.

The installation process begins. When installation finishes, a modal appears notifying you that the Ansible Automation Platform Operator is installed in the specified namespace.

- Click **View Operator** to view your newly installed Ansible Automation Platform Operator.



IMPORTANT

You can only install a single instance of the Ansible Automation Platform Operator into a single namespace. Installing multiple instances in the same namespace can lead to improper operation for both operator instances.

CHAPTER 3. INSTALLING RED HAT ANSIBLE AUTOMATION PLATFORM OPERATOR FROM THE OPENSIFT CONTAINER PLATFORM CLI


Use these instructions to install the Ansible Automation Platform Operator on Red Hat OpenShift Container Platform from the OpenShift Container Platform command-line interface (CLI) using the **oc** command.

3.1. PREREQUISITES

- Access to Red Hat OpenShift Container Platform using an account with operator installation permissions.
- The OpenShift Container Platform CLI **oc** command is installed on your local system. Refer to [Installing the OpenShift CLI](#) in the Red Hat OpenShift Container Platform product documentation for further information.

3.2. SUBSCRIBING A NAMESPACE TO AN OPERATOR USING THE OPENSIFT CONTAINER PLATFORM CLI

Use this procedure to subscribe a namespace to an operator.



IMPORTANT

You can only subscribe a single instance of the Ansible Automation Platform Operator into a single namespace. Subscribing multiple instances in the same namespace can lead to improper operation for both operator instances.

Procedure

1. Create a project for the operator.

```
oc new-project ansible-automation-platform
```

2. Create a file called **sub.yaml**.
3. Add the following YAML code to the **sub.yaml** file.

```
---
apiVersion: v1
kind: Namespace
metadata:
  labels:
    openshift.io/cluster-monitoring: "true"
  name: ansible-automation-platform
---
apiVersion: operators.coreos.com/v1
kind: OperatorGroup
metadata:
  name: ansible-automation-platform-operator
  namespace: ansible-automation-platform
spec:
```

```

targetNamespaces:
  - ansible-automation-platform
---
apiVersion: operators.coreos.com/v1alpha1
kind: Subscription
metadata:
  name: ansible-automation-platform
  namespace: ansible-automation-platform
spec:
  channel: 'stable-2.5'
  installPlanApproval: Automatic
  name: ansible-automation-platform-operator
  source: redhat-operators
  sourceNamespace: openshift-marketplace
---
apiVersion: automationcontroller.ansible.com/v1beta1
kind: AutomationController
metadata:
  name: example
  namespace: ansible-automation-platform
spec:
  replicas: 1

```

This file creates a **Subscription** object called ***ansible-automation-platform*** that subscribes the ***ansible-automation-platform*** namespace to the ***ansible-automation-platform-operator*** operator.

It then creates an **AutomationController** object called ***example*** in the ***ansible-automation-platform*** namespace.

To change the automation controller name from ***example***, edit the *name* field in the **kind: AutomationController** section of **sub.yaml** and replace ***<automation_controller_name>*** with the name you want to use:

```

apiVersion: automationcontroller.ansible.com/v1beta1
kind: AutomationController
metadata:
  name: <automation_controller_name>
  namespace: ansible-automation-platform

```

4. Run the **oc apply** command to create the objects specified in the **sub.yaml** file:

```
oc apply -f sub.yaml
```

To verify that the namespace has been successfully subscribed to the ***ansible-automation-platform-operator*** operator, run the **oc get subs** command:

```
$ oc get subs -n ansible-automation-platform
```

For further information about subscribing namespaces to operators, see [Installing from OperatorHub using the CLI](#) in the Red Hat OpenShift Container Platform *Operators* guide.

You can use the OpenShift Container Platform CLI to fetch the web address and the password of the Automation controller that you created.

3.3. FETCHING PLATFORM GATEWAY LOGIN DETAILS FROM THE OPENSIFT CONTAINER PLATFORM CLI

To login to the platform gateway, you need the web address and the password.

3.3.1. Fetching the platform gateway web address

A Red Hat OpenShift Container Platform route exposes a service at a host name, so that external clients can reach it by name. When you created the platform gateway instance, a route was created for it. The route inherits the name that you assigned to the platform gateway object in the YAML file.

Use the following command to fetch the routes:

```
oc get routes -n <platform_namespace>
```

In the following example, the **example** platform gateway is running in the **ansible-automation-platform** namespace.

```
$ oc get routes -n ansible-automation-platform
```

NAME	HOST/PORT	PATH	SERVICES	PORT	TERMINATION
WILDCARD					
example	example-ansible-automation-platform.apps-crc.testing		example-service	http	
edge/Redirect	None				

The address for the platform gateway instance is **example-ansible-automation-platform.apps-crc.testing**.

3.3.2. Fetching the platform gateway password

The YAML block for the platform gateway instance in **sub.yaml** assigns values to the *name* and *admin_user* keys. Use these values in the following command to fetch the password for the platform gateway instance.

```
oc get secret/<your instance name>-<admin_user>-password -o yaml
```

The default value for *admin_user* is **admin**. Modify the command if you changed the admin username in **sub.yaml**.

The following example retrieves the password for a platform gateway object called **example**:

```
oc get secret/example-admin-password -o yaml
```

The password for the platform gateway instance is listed in the **metadata** field in the output:

```
$ oc get secret/example-admin-password -o yaml
```

```
apiVersion: v1
data:
  password: ODzLODzLODzLODzLODzLODzLODzLODzLODzLODzL
kind: Secret
metadata:
  labels:
```

```
app.kubernetes.io/component: aap
app.kubernetes.io/name: example
app.kubernetes.io/operator-version: ""
app.kubernetes.io/part-of: example
name: example-admin-password
namespace: ansible-automation-platform
```

3.4. ADDITIONAL RESOURCES

- For more information on running operators on OpenShift Container Platform, navigate to the [OpenShift Container Platform product documentation](#) and click the *Operators - Working with Operators in OpenShift Container Platform* guide.

CHAPTER 4. CONFIGURING THE RED HAT ANSIBLE AUTOMATION PLATFORM OPERATOR ON RED HAT OPENSIFT CONTAINER PLATFORM

The platform gateway for Ansible Automation Platform enables you to manage the following Ansible Automation Platform components to form a single user interface:

- Automation controller
- Automation hub
- Event-Driven Ansible
- Red Hat Ansible Lightspeed (This feature is disabled by default, you must opt in to use it.)

Before you can deploy the platform gateway you must have Ansible Automation Platform Operator installed in a namespace. If you have not installed Ansible Automation Platform Operator see [Installing the Red Hat Ansible Automation Platform Operator on Red Hat OpenShift Container Platform](#).



NOTE

Platform gateway is only available under Ansible Automation Platform Operator version 2.5. Every component deployed under Ansible Automation Platform Operator 2.5 will also default to version 2.5.

If you have the Ansible Automation Platform Operator and some or all of the Ansible Automation Platform components installed see [Deploying the platform gateway with existing Ansible Automation Platform components](#) for how to proceed.

4.1. LINKING YOUR COMPONENTS TO THE PLATFORM GATEWAY

After installing the Ansible Automation Platform Operator in your namespace you can set up your **Ansible Automation Platform** instance. Then link all the platform components to a single user interface.

Procedure

1. Log in to Red Hat OpenShift Container Platform.
2. Navigate to **Operators → Installed Operators**.
3. Select your Ansible Automation Platform Operator deployment.
4. Select the **Details** tab.
5. On the **Ansible Automation Platform** tile click **Create instance**.
6. From the **Create Ansible Automation Platform** page enter a name for your instance in the **Name** field.
7. Click **YAML view** and paste the following:

```
spec:
  controller:
    disabled: false
```

```
eda:
  disabled: false

hub:
  disabled: false
  storage_type: file
  file_storage_storage_class: <read-write-many-storage-class>
  file_storage_size: 10Gi
```

8. Click **Create**.

Verification

Go to your Ansible Automation Platform Operator deployment and click **All instances** to verify if all instances deployed correctly. You should see the **Ansible Automation Platform** instance and the deployed **AutomationController**, **EDA**, and **AutomationHub** instances here.

Alternatively you can check by the command line, run: **oc get route**

4.2. ACCESSING THE PLATFORM GATEWAY

You should use the **Ansible Automation Platform** instance as your default. This instance links the automation controller, automation hub, and Event-Driven Ansible deployments to a single interface.

Procedure

To access your **Ansible Automation Platform** instance:

1. Log in to Red Hat OpenShift Container Platform.
2. Navigate to **Networking → Routes**
3. Click the link under **Location** for **Ansible Automation Platform**.
4. This redirects you to the Ansible Automation Platform login page. Enter "admin" as your username in the **Username** field.
5. For the password you need to:
 - a. Go to **Workloads → Secrets**.
 - b. Click **<your instance name>-admin-password** and copy the password.
 - c. Paste the password into the **Password** field.
6. Click **Login**.
7. Apply your subscription:
 - a. Click **Subscription manifest** or **Username/password**.
 - b. Upload your manifest or enter your username and password.
 - c. Select your subscription from the **Subscription** list.

- d. Click **Next**.
This redirects you to the **Analytics** page.

8. Click **Next**.
9. Select the **I agree to the terms of the license agreement** checkbox.
10. Click **Next**.

You now have access to the platform gateway user interface. If you cannot access the Ansible Automation Platform see [Section 4.4, "Frequently asked questions on platform gateway"](#) for help with troubleshooting and debugging.

4.3. DEPLOYING THE PLATFORM GATEWAY WITH EXISTING ANSIBLE AUTOMATION PLATFORM COMPONENTS

You can link any components of the Ansible Automation Platform, that you have already installed to a new **Ansible Automation Platform** instance.

The following procedure simulates a scenario where you have automation controller as an existing component and want to add automation hub and Event-Driven Ansible.

Procedure

1. Log in to Red Hat OpenShift Container Platform.
2. Navigate to **Operators → Installed Operators**.
3. Select your Ansible Automation Platform Operator deployment.
4. Click **Subscriptions** and edit your **Update channel** to **stable-2.5**.
5. Click **Details** and on the **Ansible Automation Platform** tile click **Create instance**.
6. From the **Create Ansible Automation Platform** page enter a name for your instance in the **Name** field.
7. Click **YAML view** and copy in the following:

```
apiVersion: aap.ansible.com/v1alpha1
kind: AnsibleAutomationPlatform
metadata:
  name: example-aap
  namespace: aap
spec:
  # Platform
  image_pull_policy: IfNotPresent
  # Components
  controller:
    disabled: false
    name: existing-controller-name
  eda:
    disabled: false
  hub:
    disabled: false
  ## uncomment if using file storage for Content pod
```

```

storage_type: file
file_storage_storage_class: <your-read-write-many-storage-class>
file_storage_size: 10Gi

## uncomment if using S3 storage for Content pod
# storage_type: S3
# object_storage_s3_secret: example-galaxy-object-storage


## uncomment if using Azure storage for Content pod
# storage_type: azure
# object_storage_azure_secret: azure-secret-name
lightspeed:
  disabled: true

```

- a. For new components, if you do not specify a name, a default name is generated.

8. Click **Create**.

9. To access your new instance, see [Section 4.2, “Accessing the platform gateway”](#).



NOTE

If you have an existing controller with a managed Postgres pod, after creating the **Ansible Automation Platform** resource your automation controller instance will continue to use that original Postgres pod. If you were to do a fresh install you would have a single Postgres managed pod for all instances.

4.4. FREQUENTLY ASKED QUESTIONS ON PLATFORM GATEWAY

If I delete my Ansible Automation Platform deployment will I still have access to Automation Controller?

No, automation controller, automation hub, and Event-Driven Ansible are nested within the deployment and are also deleted.

Something went wrong with my deployment but I’m not sure what, how can I find out?

You can follow along in the command line while the operator is reconciling, this can be helpful for debugging. Alternatively you can click into the deployment instance to see the status conditions being updated as the deployment goes on.

Is it still possible to view individual component logs?

When troubleshooting you should examine the **Ansible Automation Platform** instance for the main logs and then each individual component (**EDA**, **AutomationHub**, **AutomationController**) for more specific information.

Where can I view the condition of an instance?

To display status conditions click into the instance, and look under the **Details** or **Events** tab. Alternatively, to display the status conditions you can run the get command: **oc get automationcontroller <instance-name> -o jsonpath=Pipe "| jq"**

Can I track my migration in real time?

To help track the status of the migration or to understand why migration might have failed you can look at the migration logs as they are running. Use the logs command: **oc logs fresh-install-controller-migration-4.6.0-jwfm6 -f**

CHAPTER 5. CONFIGURING AUTOMATION CONTROLLER ON RED HAT OPENSIFT CONTAINER PLATFORM WEB CONSOLE

You can use these instructions to configure the automation controller operator on Red Hat OpenShift Container Platform, specify custom resources, and deploy Ansible Automation Platform with an external database.

Automation controller configuration can be done through the automation controller extra_settings or directly in the user interface after deployment. However, it is important to note that configurations made in extra_settings take precedence over settings made in the user interface.



NOTE

When an instance of automation controller is removed, the associated PVCs are not automatically deleted. This can cause issues during migration if the new deployment has the same name as the previous one. Therefore, it is recommended that you manually remove old PVCs before deploying a new automation controller instance in the same namespace. See [Finding and deleting PVCs](#) for more information.

5.1. PREREQUISITES

- You have installed the Red Hat Ansible Automation Platform catalog in Operator Hub.
- For automation controller, a default StorageClass must be configured on the cluster for the operator to dynamically create needed PVCs. This is not necessary if an external PostgreSQL database is configured.
- For Hub a StorageClass that supports ReadWriteMany must be available on the cluster to dynamically create the PVC needed for the content, redis and api pods. If it is not the default StorageClass on the cluster, you can specify it when creating your AutomationHub object.

5.1.1. Configuring your controller image pull policy

Use this procedure to configure the image pull policy on your automation controller.

Procedure

1. Log in to Red Hat OpenShift Container Platform.
2. Go to **Operators** → **Installed Operators**.
3. Select your Ansible Automation Platform Operator deployment.
4. Select the **Automation Controller** tab.
5. For new instances, click **Create AutomationController**.
 - a. For existing instances, you can edit the YAML view by clicking the **⋮** icon and then **Edit AutomationController**.
6. Click **advanced Configuration**. Under **Image Pull Policy**, click on the radio button to select
 - **Always**

- **Never**
 - **IfNotPresent**
7. To display the option under **Image Pull Secrets**, click the arrow.
 - a. Click **+** beside **Add Image Pull Secret** and enter a value.
 8. To display fields under the **Web container resource requirements** drop-down list, click the arrow.
 - a. Under **Limits**, and **Requests**, enter values for **CPU cores**, **Memory**, and **Storage**.
 9. To display fields under the **Task container resource requirements** drop-down list, click the arrow.
 - a. Under **Limits**, and **Requests**, enter values for **CPU cores**, **Memory**, and **Storage**.
 10. To display fields under the **EE Control Plane container resource requirements** drop-down list, click the arrow.
 - a. Under **Limits**, and **Requests**, enter values for **CPU cores**, **Memory**, and **Storage**.
 11. To display fields under the **PostgreSQL init container resource requirements (when using a managed service)** drop-down list, click the arrow.
 - a. Under **Limits**, and **Requests**, enter values for **CPU cores**, **Memory**, and **Storage**.
 12. To display fields under the **Redis container resource requirements** drop-down list, click the arrow.
 - a. Under **Limits**, and **Requests**, enter values for **CPU cores**, **Memory**, and **Storage**.
 13. To display fields under the **PostgreSQL container resource requirements (when using a managed instance)*** drop-down list, click the arrow.
 - a. Under **Limits**, and **Requests**, enter values for **CPU cores**, **Memory**, and **Storage**.
 14. To display the **PostgreSQL container storage requirements (when using a managed instance)** drop-down list, click the arrow.
 - a. Under **Limits**, and **Requests**, enter values for **CPU cores**, **Memory**, and **Storage**.
 15. Under **Replicas**, enter the number of instance replicas.
 16. Under **Remove used secrets on instance removal**, select **true** or **false**. The default is false.
 17. Under **Preload instance with data upon creation**, select **true** or **false**. The default is true.

5.1.2. Configuring your controller LDAP security

You can configure your LDAP SSL configuration for automation controller through any of the following options:

- The automation controller user interface.
- The platform gateway user interface. See the [Configuring LDAP authentication](#) section of the *Access management and authentication* guide for additional steps.

- The following procedure steps.

Procedure

1. If you do not have a **ldap_cacert_secret**, you can create one with the following command:

```
$ oc create secret generic <resourcename>-custom-certs \
  --from-file=ldap-ca.crt=<PATH/TO/YOUR/CA/PEM/FILE> \ 1
```

- 1 Modify this to point to where your CA cert is stored.

This will create a secret that looks like this:

```
$ oc get secret/mycerts -o yaml
apiVersion: v1
data:
  ldap-ca.crt: <mysecret> 1
kind: Secret
metadata:
  name: mycerts
  namespace: awx
type: Opaque
```

- 1 Automation controller looks for the data field **ldap-ca.crt** in the specified secret when using the **ldap_cacert_secret**.


2. Under **LDAP Certificate Authority Trust Bundle** click the drop-down menu and select your **ldap_cacert_secret**.
3. Under **LDAP Password Secret**, click the drop-down menu and select a secret.
4. Under **EE Images Pull Credentials Secret**, click the drop-down menu and select a secret.
5. Under **Bundle Cacert Secret**, click the drop-down menu and select a secret.
6. Under **Service Type**, click the drop-down menu and select
 - ClusterIP
 - LoadBalancer
 - NodePort

5.1.3. Configuring your automation controller operator route options

The Red Hat Ansible Automation Platform operator installation form allows you to further configure your automation controller operator route options under **Advanced configuration**.

Procedure


1. Log in to Red Hat OpenShift Container Platform.
2. Navigate to **Operators → Installed Operators**.

3. Select your Ansible Automation Platform Operator deployment.
4. Select the **Automation Controller** tab.
5. For new instances, click **Create AutomationController**.
 - a. For existing instances, you can edit the YAML view by clicking the  icon and then **Edit AutomationController**.
6. Click **Advanced configuration**.
7. Under **Ingress type**, click the drop-down menu and select **Route**.
8. Under **Route DNS host**, enter a common host name that the route answers to.
9. Under **Route TLS termination mechanism**, click the drop-down menu and select **Edge** or **Passthrough**. For most instances **Edge** should be selected.
10. Under **Route TLS credential secret**, click the drop-down menu and select a secret from the list.
11. Under **Enable persistence for /var/lib/projects directory** select either true or false by moving the slider.

5.1.4. Configuring the ingress type for your automation controller operator

The Ansible Automation Platform Operator installation form allows you to further configure your automation controller operator ingress under **Advanced configuration**.

Procedure

1. Log in to Red Hat OpenShift Container Platform.
2. Navigate to **Operators → Installed Operators**.
3. Select your Ansible Automation Platform Operator deployment.
4. Select the **Automation Controller** tab.
5. For new instances, click **Create AutomationController**.
 - a. For existing instances, you can edit the YAML view by clicking the  icon and then **Edit AutomationController**.
6. Click **Advanced configuration**.
7. Under **Ingress type**, click the drop-down menu and select **Ingress**.
8. Under **Ingress annotations**, enter any annotations to add to the ingress.
9. Under **Ingress TLS secret**, click the drop-down menu and select a secret from the list.

After you have configured your automation controller operator, click **Create** at the bottom of the form view. Red Hat OpenShift Container Platform will now create the pods. This may take a few minutes.

You can view the progress by navigating to **Workloads → Pods** and locating the newly created instance.

Verification

Verify that the following operator pods provided by the Ansible Automation Platform Operator installation from automation controller are running:

Operator manager controllers	Automation controller	Automation hub	Event-Driven Ansible (EDA)
<p>The operator manager controllers for each of the three operators, include the following:</p> <ul style="list-style-type: none"> • automation-controller-operator-controller-manager • automation-hub-operator-controller-manager • resource-operator-controller-manager • aap-gateway-operator-controller-manager • ansible-lightspeed-operator-controller-manager • eda-server-operator-controller-manager 	<p>After deploying automation controller, you can see the addition of the following pods:</p> <ul style="list-style-type: none"> • controller • controller-postgres • controller-web • controller-task 	<p>After deploying automation hub, you can see the addition of the following pods:</p> <ul style="list-style-type: none"> • hub-api • hub-content • hub-postgres • hub-redis • hub-worker 	<p>After deploying EDA, you can see the addition of the following pods:</p> <ul style="list-style-type: none"> • eda-activation-worker • da-api • eda-default-worker • eda-event-stream • eda-scheduler



NOTE


A missing pod can indicate the need for a pull secret. Pull secrets are required for protected or private image registries. See [Using image pull secrets](#) for more information. You can diagnose this issue further by running **oc describe pod <pod-name>** to see if there is an ImagePullBackOff error on that pod.

5.2. CONFIGURING AN EXTERNAL DATABASE FOR AUTOMATION CONTROLLER ON RED HAT ANSIBLE AUTOMATION PLATFORM OPERATOR

For users who prefer to deploy Ansible Automation Platform with an external database, they can do so by configuring a secret with instance credentials and connection information, then applying it to their cluster using the **oc create** command.

By default, the Ansible Automation Platform Operator automatically creates and configures a managed PostgreSQL pod in the same namespace as your Ansible Automation Platform deployment. You can deploy Ansible Automation Platform with an external database instead of the managed PostgreSQL pod that the Ansible Automation Platform Operator automatically creates.

Using an external database lets you share and reuse resources and manually manage backups, upgrades, and performance optimizations.



NOTE

The same external database (PostgreSQL instance) can be used for both automation hub and automation controller as long as the database names are different. In other words, you can have multiple databases with different names inside a single PostgreSQL instance.

The following section outlines the steps to configure an external database for your automation controller on a Ansible Automation Platform Operator.

Prerequisite

The external database must be a PostgreSQL database that is the version supported by the current release of Ansible Automation Platform.



NOTE

Ansible Automation Platform 2.5 supports PostgreSQL 15.

Procedure

The external postgres instance credentials and connection information must be stored in a secret, which is then set on the automation controller spec.

1. Create a **postgres_configuration_secret** YAML file, following the template below:

```
apiVersion: v1
kind: Secret
metadata:
  name: external-postgres-configuration
  namespace: <target_namespace> ❶
stringData:
  host: "<external_ip_or_url_resolvable_by_the_cluster>" ❷
  port: "<external_port>" ❸
  database: "<desired_database_name>"
  username: "<username_to_connect_as>"
  password: "<password_to_connect_with>" ❹
  sslmode: "prefer" ❺
  type: "unmanaged"
type: Opaque
```

1. Namespace to create the secret in. This should be the same namespace you want to deploy to.
 2. The resolvable hostname for your database node.
 3. External port defaults to **5432**.
 4. Value for variable **password** should not contain single or double quotes (' ') or backslashes (\) to avoid any issues during deployment, backup or restoration.
 5. The variable **sslmode** is valid for **external** databases only. The allowed values are: **prefer**, **disable**, **allow**, **require**, **verify-ca**, and **verify-full**.
2. Apply **external-postgres-configuration-secret.yml** to your cluster using the **oc create** command.

```
$ oc create -f external-postgres-configuration-secret.yml
```

3. When creating your **AutomationController** custom resource object, specify the secret on your spec, following the example below:

```
apiVersion: automationcontroller.ansible.com/v1beta1
kind: AutomationController
metadata:
  name: controller-dev
spec:
  postgres_configuration_secret: external-postgres-configuration
```

5.3. FINDING AND DELETING PVCS

A persistent volume claim (PVC) is a storage volume used to store data that automation hub and automation controller applications use. These PVCs are independent from the applications and remain even when the application is deleted. If you are confident that you no longer need a PVC, or have backed it up elsewhere, you can manually delete them.

Procedure

1. List the existing PVCs in your deployment namespace:

```
oc get pvc -n <namespace>
```

2. Identify the PVC associated with your previous deployment by comparing the old deployment name and the PVC name.
3. Delete the old PVC:

```
oc delete pvc -n <namespace> <pvc-name>
```

5.4. ADDITIONAL RESOURCES

- For more information on running operators on OpenShift Container Platform, navigate to the [OpenShift Container Platform product documentation](#) and click the *Operators - Working with Operators in OpenShift Container Platform* guide.

CHAPTER 6. CONFIGURING AUTOMATION HUB ON RED HAT OPENSIFT CONTAINER PLATFORM WEB CONSOLE

You can use these instructions to configure the automation hub operator on Red Hat OpenShift Container Platform, specify custom resources, and deploy Ansible Automation Platform with an external database.

Automation hub configuration can be done through the automation hub pulp_settings or directly in the user interface after deployment. However, it is important to note that configurations made in pulp_settings take precedence over settings made in the user interface. Hub settings should always be set as lowercase on the Hub custom resource specification.



NOTE

When an instance of automation hub is removed, the PVCs are not automatically deleted. This can cause issues during migration if the new deployment has the same name as the previous one. Therefore, it is recommended that you manually remove old PVCs before deploying a new automation hub instance in the same namespace. See [Finding and deleting PVCs](#) for more information.

6.1. PREREQUISITES

- You have installed the Ansible Automation Platform Operator in Operator Hub.

6.1.1. Storage options for Ansible Automation Platform Operator installation on Red Hat OpenShift Container Platform

Automation hub requires **ReadWriteMany** file-based storage, Azure Blob storage, or Amazon S3-compliant storage for operation so that multiple pods can access shared content, such as collections.

The process for configuring object storage on the **AutomationHub** CR is similar for Amazon S3 and Azure Blob Storage.

If you are using file-based storage and your installation scenario includes automation hub, ensure that the storage option for Ansible Automation Platform Operator is set to **ReadWriteMany**. **ReadWriteMany** is the default storage option.

In addition, OpenShift Data Foundation provides a **ReadWriteMany** or S3-compliant implementation. Also, you can set up NFS storage configuration to support **ReadWriteMany**. This, however, introduces the NFS server as a potential, single point of failure.

Additional resources

- [Persistent storage using NFS](#) in the OpenShift Container Platform *Storage* guide
- IBM's [How do I create a storage class for NFS dynamic storage provisioning in an OpenShift environment?](#)

6.1.1.1. Provisioning OCP storage with **ReadWriteMany** access mode

To ensure successful installation of Ansible Automation Platform Operator, you must provision your storage type for automation hub initially to **ReadWriteMany** access mode.

Procedure

1. Go to **Storage** → **PersistentVolume**.
2. Click btn: Create PersistentVolume.
3. In the first step, update the **accessModes** from the default **ReadWriteOnce** to **ReadWriteMany**.
 - a. See [Provisioning](#) to update the access mode. for a detailed overview.
4. Complete the additional steps in this section to create the persistent volume claim (PVC).

6.1.1.2. Configuring object storage on Amazon S3

Red Hat supports Amazon Simple Storage Service (S3) for automation hub. You can configure it when deploying the **AutomationHub** custom resource (CR), or you can configure it for an existing instance.

Prerequisites

- Create an Amazon S3 bucket to store the objects.
- Note the name of the S3 bucket.

Procedure

1. Create a Kubernetes secret containing the AWS credentials and connection details, and the name of your Amazon S3 bucket. The following example creates a secret called **test-s3**:

```
$ oc -n $HUB_NAMESPACE apply -f- <<EOF
apiVersion: v1
kind: Secret
metadata:
  name: 'test-s3'
stringData:
  s3-access-key-id: $S3_ACCESS_KEY_ID
  s3-secret-access-key: $S3_SECRET_ACCESS_KEY
  s3-bucket-name: $S3_BUCKET_NAME
  s3-region: $S3_REGION
EOF
```

2. Add the secret to the automation hub custom resource (CR) **spec**:

```
spec:
  object_storage_s3_secret: test-s3
```

3. If you are applying this secret to an existing instance, restart the API pods for the change to take effect. **<hub-name>** is the name of your hub instance.

```
$ oc -n $HUB_NAMESPACE delete pod -l app.kubernetes.io/name=<hub-name>-api
```

6.1.1.3. Configuring object storage on Azure Blob

Red Hat supports Azure Blob Storage for automation hub. You can configure it when deploying the **AutomationHub** custom resource (CR), or you can configure it for an existing instance.

Prerequisites

- Create an Azure Storage blob container to store the objects.
- Note the name of the blob container.

Procedure

1. Create a Kubernetes secret containing the credentials and connection details for your Azure account, and the name of your Azure Storage blob container. The following example creates a secret called **test-azure**:

```
$ oc -n $HUB_NAMESPACE apply -f -<EOF
apiVersion: v1
kind: Secret
metadata:
  name: 'test-azure'
stringData:
  azure-account-name: $AZURE_ACCOUNT_NAME
  azure-account-key: $AZURE_ACCOUNT_KEY
  azure-container: $AZURE_CONTAINER
  azure-container-path: $AZURE_CONTAINER_PATH
  azure-connection-string: $AZURE_CONNECTION_STRING
EOF
```

2. Add the secret to the automation hub custom resource (CR) **spec**:

```
spec:
  object_storage_azure_secret: test-azure
```

3. If you are applying this secret to an existing instance, restart the API pods for the change to take effect. **<hub-name>** is the name of your hub instance.


```
$ oc -n $HUB_NAMESPACE delete pod -l app.kubernetes.io/name=<hub-name>-api
```

6.1.2. Configure your automation hub operator route options

The Red Hat Ansible Automation Platform operator installation form allows you to further configure your automation hub operator route options under **Advanced configuration**.

Procedure


1. Log in to Red Hat OpenShift Container Platform.
2. Navigate to **Operators → Installed Operators**.
3. Select your Ansible Automation Platform Operator deployment.
4. Select the **Automation Hub** tab.
5. For new instances, click **Create AutomationHub**.

- a. For existing instances, you can edit the YAML view by clicking the  icon and then **Edit AutomationHub**.
6. Click **Advanced configuration**.
7. Under **Ingress type**, click the drop-down menu and select **Route**.
8. Under **Route DNS host**, enter a common host name that the route answers to.
9. Under **Route TLS termination mechanism**, click the drop-down menu and select **Edge** or **Passthrough**.
10. Under **Route TLS credential secret**, click the drop-down menu and select a secret from the list.

6.1.3. Configuring the ingress type for your automation hub operator

The Ansible Automation Platform Operator installation form allows you to further configure your automation hub operator ingress under **Advanced configuration**.

Procedure

1. Log in to Red Hat OpenShift Container Platform.
2. Navigate to **Operators → Installed Operators**.
3. Select your Ansible Automation Platform Operator deployment.
4. Select the **Automation Hub** tab.
5. For new instances, click **Create AutomationHub**.
 - a. For existing instances, you can edit the YAML view by clicking the  icon and then **Edit AutomationHub**.
6. Click **Advanced Configuration**.
7. Under **Ingress type**, click the drop-down menu and select **Ingress**.
8. Under **Ingress annotations**, enter any annotations to add to the ingress.
9. Under **Ingress TLS secret**, click the drop-down menu and select a secret from the list.

After you have configured your automation hub operator, click **Create** at the bottom of the form view. Red Hat OpenShift Container Platform will now create the pods. This may take a few minutes.

You can view the progress by navigating to **Workloads → Pods** and locating the newly created instance.

Verification

Verify that the following operator pods provided by the Ansible Automation Platform Operator installation from automation hub are running:

Operator manager controllers	automation controller	automation hub
<p>The operator manager controllers for each of the 3 operators, include the following:</p> <ul style="list-style-type: none"> • automation-controller-operator-controller-manager • automation-hub-operator-controller-manager • resource-operator-controller-manager 	<p>After deploying automation controller, you will see the addition of these pods:</p> <ul style="list-style-type: none"> • controller • controller-postgres 	<p>After deploying automation hub, you will see the addition of these pods:</p> <ul style="list-style-type: none"> • hub-api • hub-content • hub-postgres • hub-redis • hub-worker



NOTE

A missing pod can indicate the need for a pull secret. Pull secrets are required for protected or private image registries. See [Using image pull secrets](#) for more information. You can diagnose this issue further by running `oc describe pod <pod-name>` to see if there is an ImagePullBackOff error on that pod.

6.2. CONFIGURING LDAP AUTHENTICATION FOR ANSIBLE AUTOMATION HUB ON OPENSIFT CONTAINER PLATFORM

Configure LDAP authentication settings for Ansible Automation Platform on OpenShift Container Platform in the spec section of your Hub instance configuration file.

Procedure

- Use the following example to configure LDAP in your automation hub instance. For any blank fields, enter ``.

```
spec:
  pulp_settings:
    auth_ldap_user_attr_map:
      email: "mail"
      first_name: "givenName"
      last_name: "sn"
    auth_ldap_group_search_base_dn: 'cn=groups,cn=accounts,dc=example,dc=com'
    auth_ldap_bind_dn: ''
    auth_ldap_bind_password: ''
    auth_ldap_group_search_filter: (objectClass=posixGroup)
    auth_ldap_user_search_scope: SUBTREE
    auth_ldap_server_uri: 'ldap://ldapserver:389'
    authentication_backend_preset: ldap
    auth_ldap_mirror_groups: 'True'
    auth_ldap_user_search_base_dn: 'cn=users,cn=accounts,dc=example,dc=com'
    auth_ldap_bind_password: 'ldappassword'
```

```
auth_ldap_user_search_filter: (uid=%(user)s)
auth_ldap_group_search_scope: SUBTREE
auth_ldap_user_flags_by_group: '@json {"is_superuser": "cn=tower-
admin,cn=groups,cn=accounts,dc=example,dc=com"}'
```

**NOTE**

Do not leave any fields empty. For fields with no variable, enter `` to indicate a default value.

6.3. FINDING THE AUTOMATION HUB ROUTE

You can access the automation hub through the platform gateway or through the following procedure.

Procedure

1. Log into Red Hat OpenShift Container Platform.
2. Navigate to **Networking → Routes**.
3. Under **Location**, click on the URL for your automation hub instance.

The automation hub user interface launches where you can sign in with the administrator credentials specified during the operator configuration process.

**NOTE**

If you did not specify an administrator password during configuration, one was automatically created for you. To locate this password, go to your project, select **Workloads → Secrets** and open controller-admin-password. From there you can copy the password and paste it into the Automation hub password field.

6.4. CONFIGURING AN EXTERNAL DATABASE FOR AUTOMATION HUB ON RED HAT ANSIBLE AUTOMATION PLATFORM OPERATOR

For users who prefer to deploy Ansible Automation Platform with an external database, they can do so by configuring a secret with instance credentials and connection information, then applying it to their cluster using the **oc create** command.

By default, the Ansible Automation Platform Operator automatically creates and configures a managed PostgreSQL pod in the same namespace as your Ansible Automation Platform deployment.

You can choose to use an external database instead if you prefer to use a dedicated node to ensure dedicated resources or to manually manage backups, upgrades, or performance tweaks.


**NOTE**

The same external database (PostgreSQL instance) can be used for both automation hub and automation controller as long as the database names are different. In other words, you can have multiple databases with different names inside a single PostgreSQL instance.

The following section outlines the steps to configure an external database for your automation hub on a Ansible Automation Platform Operator.

Prerequisite

The external database must be a PostgreSQL database that is the version supported by the current release of Ansible Automation Platform.



NOTE

Ansible Automation Platform 2.5 supports PostgreSQL 15.

Procedure

The external postgres instance credentials and connection information will need to be stored in a secret, which will then be set on the automation hub spec.

1. Create a **postgres_configuration_secret** YAML file, following the template below:

```
apiVersion: v1
kind: Secret
metadata:
  name: external-postgres-configuration
  namespace: <target_namespace> ❶
stringData:
  host: "<external_ip_or_url_resolvable_by_the_cluster>" ❷
  port: "<external_port>" ❸
  database: "<desired_database_name>"
  username: "<username_to_connect_as>"
  password: "<password_to_connect_with>" ❹
  sslmode: "prefer" ❺
  type: "unmanaged"
type: Opaque
```

- ❶ Namespace to create the secret in. This should be the same namespace you want to deploy to.
- ❷ The resolvable hostname for your database node.
- ❸ External port defaults to **5432**.
- ❹ Value for variable **password** should not contain single or double quotes (' ') or backslashes (\) to avoid any issues during deployment, backup or restoration.
- ❺ The variable **sslmode** is valid for **external** databases only. The allowed values are: **prefer**, **disable**, **allow**, **require**, **verify-ca**, and **verify-full**.

2. Apply **external-postgres-configuration-secret.yml** to your cluster using the **oc create** command.

```
$ oc create -f external-postgres-configuration-secret.yml
```

- When creating your **AutomationHub** custom resource object, specify the secret on your spec, following the example below:

```
apiVersion: automationhub.ansible.com/v1beta1
kind: AutomationHub
metadata:
  name: hub-dev
spec:
  postgres_configuration_secret: external-postgres-configuration
```

6.4.1. Enabling the hstore extension for the automation hub PostgreSQL database

Added in Ansible Automation Platform 2.5, the database migration script uses **hstore** fields to store information, therefore the **hstore** extension to the automation hub PostgreSQL database must be enabled.

This process is automatic when using the Ansible Automation Platform installer and a managed PostgreSQL server.

If the PostgreSQL database is external, you must enable the **hstore** extension to the automation hub PostgreSQL database manually before automation hub installation.

If the **hstore** extension is not enabled before automation hub installation, a failure is raised during database migration.

Procedure

- Check if the extension is available on the PostgreSQL server (automation hub database).

```
$ psql -d <automation hub database> -c "SELECT * FROM pg_available_extensions WHERE
name='hstore'"
```

- Where the default value for **<automation hub database>** is **automationhub**.

Example output with hstore available:

```
name | default_version | installed_version | comment
-----+-----+-----+-----
hstore | 1.7             |                  | data type for storing sets of (key, value) pairs
(1 row)
```

Example output with hstore not available:

```
name | default_version | installed_version | comment
-----+-----+-----+-----
(0 rows)
```

- On a RHEL based server, the **hstore** extension is included in the **postgresql-contrib** RPM package, which is not installed automatically when installing the PostgreSQL server RPM package.

To install the RPM package, use the following command:

```
dnf install postgresql-contrib
```

4. Create the **hstore** PostgreSQL extension on the automation hub database with the following command:

```
$ psql -d <automation hub database> -c "CREATE EXTENSION hstore;"
```

The output of which is:

```
CREATE EXTENSION
```

In the following output, the **installed_version** field contains the **hstore** extension used, indicating that **hstore** is enabled.

```
name | default_version | installed_version | comment
-----+-----+-----+-----
hstore | 1.7 | 1.7 | data type for storing sets of (key, value) pairs
(1 row)
```

6.5. FINDING AND DELETING PVCS

A persistent volume claim (PVC) is a storage volume used to store data that automation hub and automation controller applications use. These PVCs are independent from the applications and remain even when the application is deleted. If you are confident that you no longer need a PVC, or have backed it up elsewhere, you can manually delete them.

Procedure

1. List the existing PVCs in your deployment namespace:

```
oc get pvc -n <namespace>
```

2. Identify the PVC associated with your previous deployment by comparing the old deployment name and the PVC name.
3. Delete the old PVC:

```
oc delete pvc -n <namespace> <pvc-name>
```

6.6. ADDITIONAL CONFIGURATIONS

A collection download count can help you understand collection usage. To add a collection download count to automation hub, set the following configuration:

```
spec:
  pulp_settings:
    ansible_collect_download_count: true
```

When **ansible_collect_download_count** is enabled, automation hub will display a download count by the collection.

6.7. ADDITIONAL RESOURCES

- For more information on running operators on OpenShift Container Platform, navigate to the [OpenShift Container Platform product documentation](#) and click the *Operators - Working with Operators in OpenShift Container Platform* guide.

CHAPTER 7. DEPLOYING CLUSTERED REDIS ON RED HAT ANSIBLE AUTOMATION PLATFORM OPERATOR


When you create an Ansible Automation Platform instance through the Ansible Automation Platform Operator, standalone Redis is assigned by default. To deploy clustered Redis, use the following procedure.

For more information about Redis, refer to Caching and queueing system in the *Planning your installation* guide.

Prerequisites

- You have installed an Ansible Automation Platform Operator deployment.

Procedure

1. Log in to Red Hat OpenShift Container Platform.
2. Navigate to **Operators** → **Installed Operators**.
3. Select your Ansible Automation Platform Operator deployment.
4. Select the **Details** tab.
5. On the **Ansible Automation Platform** tile click **Create instance**.
 - a. For existing instances, you can edit the YAML view by clicking the  icon and then **Edit AnsibleAutomationPlatform**.
 - b. Change the **redis_mode** value to "clustered".
 - c. Click **Reload**, then **Save**.
6. Click to expand **Advanced configuration**.
7. For the **Redis Mode** list, select **Cluster**.
8. Configure the rest of your instance as necessary, then click **Create**.

Your instance will deploy with a cluster Redis with 6 Redis replicas as default.

CHAPTER 8. USING RED HAT SINGLE SIGN-ON OPERATOR WITH AUTOMATION HUB

Private automation hub uses Red Hat Single Sign-On for authentication.

The Red Hat Single Sign-On Operator creates and manages resources. Use this operator to create custom resources to automate Red Hat Single Sign-On administration in OpenShift.

- When installing Ansible Automation Platform on *Virtual Machines* (VMs) the installer can automatically install and configure Red Hat Single Sign-On for use with private automation hub.
- When installing Ansible Automation Platform on Red Hat OpenShift Container Platform you must install Single Sign-On separately.

This chapter describes the process to configure Red Hat Single Sign-On and integrate it with private automation hub when Ansible Automation Platform is installed on OpenShift Container Platform.

Prerequisites

- You have access to Red Hat OpenShift Container Platform using an account with operator installation permissions.
- You have installed the catalog containing the Red Hat Ansible Automation Platform operators.
- You have installed the Red Hat Single Sign-On Operator. To install the Red Hat Single Sign-On Operator, follow the procedure in [Installing Red Hat Single Sign-On using a custom resource](#) in the Red Hat Single Sign-On documentation.

8.1. CREATING A KEYCLOAK INSTANCE

After you install the Red Hat Single Sign-On Operator, you can create a Keycloak instance for use with Ansible Automation Platform.

From here you provide an external Postgres or one will be created for you.

Procedure

/ Log in to Red Hat OpenShift Container Platform. . Navigate to **Operator → Installed Operators**. . Select the **RH-SSO** project. . Select the **Red Hat Single Sign-On Operator**. . On the Red Hat Single Sign-On Operator details page select **Keycloak**. . Click **Create instance**. . Click **YAML view**.

+ The default Keycloak custom resource is as follows:

+

```
apiVersion: keycloak.org/v1alpha1
kind: Keycloak
metadata:
  name: example-keycloak
labels:
  app: sso
  namespace: aap
spec:
```



```
externalAccess:
enabled: true
instances: 1
```

+ . Click **Create**.

1. When deployment is complete, you can use this credential to login to the administrative console.
2. You can find the credentials for the administrator in the **credential-<custom-resource>** (example keycloak) secret in the namespace.

8.2. CREATING A KEYCLOAK REALM FOR ANSIBLE AUTOMATION PLATFORM

Create a realm to manage a set of users, credentials, roles, and groups. A user belongs to and logs into a realm. Realms are isolated from one another and can only manage and authenticate the users that they control.

Procedure

1. Log in to Red Hat OpenShift Container Platform.
2. Navigate to **Operator → Installed Operators**.
3. Select the **Red Hat Single Sign-On Operator** project.
4. Select the **Keycloak Realm** tab and click **Create Keycloak Realm**.
5. On the **Keycloak Realm** form, select **YAML view**. Edit the YAML file as follows:

```
kind: KeycloakRealm
apiVersion: keycloak.org/v1alpha1
metadata:
  name: ansible-automation-platform-keycloakrealm
  namespace: rh-sso
  labels:
    app: sso
    realm: ansible-automation-platform
spec:
  realm:
    id: ansible-automation-platform
    realm: ansible-automation-platform
    enabled: true
    displayName: Ansible Automation Platform
  instanceSelector:
    matchLabels:
      app: sso
```

Field	Description
metadata.name	Set a unique value in metadata for the name of the configuration resource (CR).

metadata.namespace	Set a unique value in metadata for the name of the configuration resource (CR).
metadata.labels.app	Set labels to a unique value. This is used when creating the client CR.
metadata.labels.realm	Set labels to a unique value. This is used when creating the client CR.
spec.realm.id	Set the realm name and id. These must be the same.
spec.realm.realm	Set the realm name and id. These must be the same.
spec.realm.displayname	Set the name to display.

- Click **Create** and wait for the process to complete.

8.3. CREATING A KEYCLOAK CLIENT

Keycloak clients authenticate hub users with Red Hat Single Sign-On. When a user authenticates the request goes through the Keycloak client. When Single Sign-On validates or issues the **OAuth** token, the client provides the response to automation hub and the user can log in.

Procedure

- Log in to Red Hat OpenShift Container Platform.
- Navigate to **Operator → Installed Operators**.
- Select the Red Hat Single Sign-On Operator project.
- Select the **Keycloak Client** tab and click **Create Keycloak Client**.
- On the Keycloak Realm form, select **YAML view**.
- Replace the default YAML file with the following:

```
kind: KeycloakClient
apiVersion: keycloak.org/v1alpha1
metadata:
  name: automation-hub-client-secret
  labels:
    app: sso
    realm: ansible-automation-platform
  namespace: rh-sso
spec:
  realmSelector:
    matchLabels:
      app: sso
      realm: ansible-automation-platform
```

```

client:
  name: Automation Hub
  clientId: automation-hub
  secret: <client-secret>
  clientAuthenticatorType: client-secret
  description: Client for automation hub
  attributes:
    user.info.response.signature.alg: RS256
    request.object.signature.alg: RS256
  directAccessGrantsEnabled: true
  publicClient: true
  protocol: openid-connect
  standardFlowEnabled: true
  protocolMappers:
    - config:
        access.token.claim: "true"
        claim.name: "family_name"
        id.token.claim: "true"
        jsonType.label: String
        user.attribute: lastName
        userinfo.token.claim: "true"
      consentRequired: false
      name: family name
      protocol: openid-connect
      protocolMapper: oidc-usermodel-property-mapper
    - config:
        userinfo.token.claim: "true"
        user.attribute: email
        id.token.claim: "true"
        access.token.claim: "true"
        claim.name: email
        jsonType.label: String
        name: email
        protocol: openid-connect
        protocolMapper: oidc-usermodel-property-mapper
        consentRequired: false
    - config:
        multivalued: "true"
        access.token.claim: "true"
        claim.name: "resource_access.${client_id}.roles"
        jsonType.label: String
        name: client roles
        protocol: openid-connect
        protocolMapper: oidc-usermodel-client-role-mapper
        consentRequired: false
    - config:
        userinfo.token.claim: "true"
        user.attribute: firstName
        id.token.claim: "true"
        access.token.claim: "true"
        claim.name: given_name
        jsonType.label: String
        name: given name
        protocol: openid-connect
        protocolMapper: oidc-usermodel-property-mapper
        consentRequired: false

```

```

- config:
  id.token.claim: "true"
  access.token.claim: "true"
  userinfo.token.claim: "true"
  name: full name
  protocol: openid-connect
  protocolMapper: oidc-full-name-mapper
  consentRequired: false
- config:
  userinfo.token.claim: "true"
  user.attribute: username
  id.token.claim: "true"
  access.token.claim: "true"
  claim.name: preferred_username
  jsonType.label: String
  name: <username>
  protocol: openid-connect
  protocolMapper: oidc-usermodel-property-mapper
  consentRequired: false
- config:
  access.token.claim: "true"
  claim.name: "group"
  full.path: "true"
  id.token.claim: "true"
  userinfo.token.claim: "true"
  consentRequired: false
  name: group
  protocol: openid-connect
  protocolMapper: oidc-group-membership-mapper
- config:
  multivalued: 'true'
  id.token.claim: 'true'
  access.token.claim: 'true'
  userinfo.token.claim: 'true'
  usermodel.clientRoleMapping.clientId: 'automation-hub'
  claim.name: client_roles
  jsonType.label: String
  name: client_roles
  protocolMapper: oidc-usermodel-client-role-mapper
  protocol: openid-connect
- config:
  id.token.claim: "true"
  access.token.claim: "true"
  included.client.audience: 'automation-hub'
  protocol: openid-connect
  name: audience mapper
  protocolMapper: oidc-audience-mapper
roles:
- name: "hubadmin"
  description: "An administrator role for automation hub"

```

1

Replace this with a unique value.

7. Click **Create** and wait for the process to complete.

After you deploy automation hub, you must update the client with the “Valid Redirect URIs” and “Web Origins” as described in [Updating the Red Hat Single Sign-On client](#). Additionally, the client comes pre-configured with token mappers, however, if your authentication provider does not provide group data to Red Hat SSO, then the group mapping must be updated to reflect how that information is passed. This is commonly by user attribute.

8.4. CREATING A KEYCLOAK USER

This procedure creates a Keycloak user, with the **hubadmin** role, that can log in to automation hub with Super Administration privileges.

Procedure

1. Log in to Red Hat OpenShift Container Platform.
2. Navigate to **Operator → Installed Operators**.
3. Select the Red Hat Single Sign-On Operator project.
4. Select the **Keycloak Realm** tab and click **Create Keycloak User**.
5. On the **Keycloak User** form, select **YAML view**.
6. Replace the default YAML file with the following:

```
apiVersion: keycloak.org/v1alpha1
kind: KeycloakUser
metadata:
  name: hubadmin-user
  labels:
    app: sso
    realm: ansible-automation-platform
  namespace: rh-sso
spec:
  realmSelector:
    matchLabels:
      app: sso
      realm: ansible-automation-platform
  user:
    username: hub_admin
    firstName: Hub
    lastName: Admin
    email: hub_admin@example.com
    enabled: true
    emailVerified: false
    credentials:
      - type: password
        value: <ch8ngeme>
    clientRoles:
      automation-hub:
        - hubadmin
```

7. Click **Create** and wait for the process to complete.

When a user is created, the Operator creates a Secret containing both the username and password

using the following naming pattern: **credential-`<realm name>`-`<username>`-`<namespace>`**. In this example the credential is called **credential-ansible-automation-platform-hub-admin-rh-sso**. When a user is created the Operator does not update the user's password. Password changes are not reflected in the Secret.

8.5. INSTALLING THE ANSIBLE AUTOMATION PLATFORM OPERATOR

Procedure

1. Log in to Red Hat OpenShift Container Platform.
2. Navigate to **Operator → Operator Hub**.
3. Search for the Ansible Automation Platform Operator.
4. Select the Ansible Automation Platform Operator project.
5. Click on the Operator tile.
6. Click **Install**.
7. Select a Project to install the Operator into. Red Hat recommends using the Operator recommended Namespace name.
 - a. If you want to install the Operator into a project other than the recommended one, select **Create Project** from the drop down menu.
 - b. Enter the Project name.
 - c. Click **Create**.
8. Click **Install**.
9. When the Operator has been installed, click **View Operator**.

8.6. CREATING A RED HAT SINGLE SIGN-ON CONNECTION SECRET

Use this procedure to create a connection secret for Red Hat Single Sign-On.

Procedure

1. Navigate to **`https://<sso_host>/auth/realms/ansible-automation-platform`**.
2. Copy the **public_key** value.
3. In the OpenShift Web UI, navigate to **Workloads → Secrets**.
4. Select the **ansible-automation-platform** project.
5. Click **Create**, and select **From YAML**.
6. Edit the following YAML to create the secret

```
apiVersion: v1
kind: Secret
```

```

metadata:
  name: automation-hub-sso
  namespace: ansible-automation-platform
type: Opaque
stringData:
  keycloak_host: "keycloak-rh-sso.apps-crc.testing"
  keycloak_port: "443"
  keycloak_protocol: "https"
  keycloak_realm: "ansible-automation-platform"
  keycloak_admin_role: "hubadmin"
  social_auth_keycloak_key: "automation-hub"
  social_auth_keycloak_secret: "client-secret"
  social_auth_keycloak_public_key: >-

```

- 1 This name is used in the next step when creating the automation hub instance.
- 2 If the secret was changed when creating the Keycloak client for automation hub be sure to change this value to match.
- 3 Enter the value of the **public_key** copied in [Installing the Ansible Automation Platform Operator](#).

7. Click **Create** and wait for the process to complete.

8.7. INSTALLING AUTOMATION HUB USING THE ANSIBLE AUTOMATION PLATFORM OPERATOR

Use the following procedure to install automation hub using the Ansible Automation Platform Operator.

Procedure

1. Log in to Red Hat OpenShift Container Platform.
2. Navigate to **Operator** → **Installed Operators**.
3. Select your Ansible Automation Platform Operator deployment.
4. Select the Automation hub tab.
5. Click **Create Automation hub**.
6. Select **YAML view**. The YAML should be similar to:

```

apiVersion: automationhub.ansible.com/v1beta1
kind: AutomationHub
metadata:
  name: private-ah
  namespace: aap
spec:
  sso_secret: automation-hub-sso
  pulp_settings:
    verify_ssl: false
  route_tls_termination_mechanism: Edge
  ingress_type: Route

```

```

loadbalancer_port: 80
file_storage_size: 100Gi
image_pull_policy: IfNotPresent
replicas: 1
web_replicas: N
task_replicas: N
file_storage_access_mode: ReadWriteMany
content:
  log_level: INFO
  replicas: 2
postgres_storage_requirements:
  limits:
    storage: 50Gi
  requests:
    storage: 8Gi
api:
  log_level: INFO
  replicas: 1
postgres_resource_requirements:
  limits:
    cpu: 1000m
    memory: 8Gi
  requests:
    cpu: 500m
    memory: 2Gi
loadbalancer_protocol: http
resource_manager:
  replicas: 1
worker:
  replicas: 2

```

- 1 Set `metadata.name` to the name to use for the instance.
- 2 Set `spec.sso_secret` to the name of the secret created in [Creating a Secret to hold the Red Hat Single Sign On connection details](#).
- 3 Scale replicas up or down for each deployment by using the **`web_replicas`** or **`task_replicas`** respectively, where N represents the number of replicas you want to create. Alternatively, you can scale all pods across both deployments by using **`replicas`**. See [Scaling the Web and Task Pods independently](#) for details.



NOTE

This YAML turns off SSL verification (**`ssl_verify: false`**). If you are not using self-signed certificates for OpenShift this setting can be removed.

7. Click **Create** and wait for the process to complete.

8.8. DETERMINING THE AUTOMATION HUB ROUTE

Use the following procedure to determine the hub route.

Procedure

1. Log in to Red Hat OpenShift Container Platform.
2. Navigate to **Networking → Routes**.
3. Select the project you used for the install.
4. Copy the location of the **private-ah-web-svc** service. The name of the service is different if you used a different name when creating the automation hub instance. This is used later to update the Red Hat Single Sign-On client.

8.9. UPDATING THE RED HAT SINGLE SIGN-ON CLIENT


After you install automation hub and you know the URL of the instance, you must update the Red Hat Single Sign-On to set the Valid Redirect URIs and Web Origins settings.

Procedure

1. Log in to Red Hat OpenShift Container Platform.
2. Navigate to **Operator → Installed Operators**.
3. Select the **RH-SSO** project.
4. Click **Red Hat Single Sign-On Operator**.
5. Select **Keycloak Client**.
6. Click on the automation-hub-client-secret client.
7. Select **YAML**.
8. Update the Client YAML to add the Valid Redirect URIs and Web Origins settings.

```
redirectUris:
  - 'https://private-ah-ansible-automation-platform.apps-crc.testing/*'
webOrigins:
  - 'https://private-ah-ansible-automation-platform.apps-crc.testing'
```

Field	Description
redirectURIs	This is the location determined in Determine Automation Hub Route . Be sure to add the <code>/*</code> to the end of the redirectUris setting.
webOrigins	This is the location determined in Determine Automation Hub Route .



NOTE

Ensure the indentation is correct when entering these settings.

9. Click **Save**.

To verify connectivity

1. Navigate to the automation hub route.
2. Enter the **hub_admin** user credentials and sign in.
3. Red Hat Single Sign-On processes the authentication and redirects back to automation hub.

8.10. ADDITIONAL RESOURCES

- For more information on running operators on OpenShift Container Platform, see [Working with Operators in OpenShift Container Platform](#) in the OpenShift Container Platform product documentation.


CHAPTER 9. ADDING EXECUTION NODES TO RED HAT ANSIBLE AUTOMATION PLATFORM OPERATOR

You can enable the Ansible Automation Platform Operator with execution nodes by downloading and installing the install bundle.

Prerequisites

- An automation controller instance.
- The receptor collection package is installed.
- AAP Repository **ansible-automation-platform-2.4-for-rhel-{RHEL-RELEASE-NUMBER}-x86_64-rpms** is enabled.

Procedure

1. Log in to Red Hat Ansible Automation Platform.
2. In the navigation panel, select **Automation Execution** → **Infrastructure** → **Instances**.
3. Click **Add**.
4. Input the Execution Node domain name or IP in the **Host Name** field.
5. Optional: Input the port number in the **Listener Port** field.
6. Click **Save**.
7. Click the download icon  next to **Install Bundle**. This starts a download, take note of where you save the file
8. Untar the gz file.



NOTE

To run the **install_receptor.yml** playbook you need to install the receptor collection from Ansible Galaxy: **Ansible-galaxy collection install -r requirements.yml**

9. Update the playbook with your user name and SSH private key file. Note that **ansible_host** pre-populates with the hostname you input earlier.

```
all:
  hosts:
    remote-execution:
      ansible_host: example_host_name # Same with configured in AAP WebUI
      ansible_user: <username> #user provided
      Ansible_ssh_private_key_file: ~/.ssh/id_example
```

10. Open your terminal, and navigate to the directory where you saved the playbook.
11. To install the bundle run:

```
ansible-playbook install_receptor.yml -i inventory.yml
```

12. When installed you can now upgrade your execution node by downloading and re-running the playbook for the instance you created.

Verification

To verify receptor service status run the following command:

```
sudo systemctl status receptor.service
```

Make sure the service is in **active (running)** state

To verify if your playbook runs correctly on your new node run the following command:

```
watch podman ps
```

Additional resources

- For more information about managing instance groups see the [Managing Instance Groups](#) section of the Automation Controller User Guide.

CHAPTER 10. ANSIBLE AUTOMATION PLATFORM RESOURCE OPERATOR

10.1. RESOURCE OPERATOR OVERVIEW

Resource Operator is a custom resource (CR) that you can deploy after you have created your platform gateway deployment. With Resource Operator you can define projects, job templates, and inventories through the use of YAML files. These YAML files are then used by automation controller to create these resources. You can create the YAML through the **Form view** that prompts you for keys and values for your YAML code. Alternatively, to work with YAML directly, you can select **YAML view**.

There are currently two custom resources provided by the Resource Operator:

- **AnsibleJob**: launches a job in the automation controller instance specified in the Kubernetes secret (automation controller host URL, token).
- **JobTemplate**: creates a job template in the automation controller instance specified.

10.2. USING RESOURCE OPERATOR

The Resource Operator itself does not do anything until the user creates an object. As soon as the user creates an **AutomationControllerProject** or **AnsibleJob** resource, the Resource Operator will start processing that object.

Prerequisites

- Install the Kubernetes-based cluster of your choice.
- Deploy automation controller using the **automation-controller-operator**.

After installing the **automation-controller-resource-operator** in your cluster, you must create a Kubernetes (k8s) secret with the connection information for your automation controller instance. Then you can use Resource Operator to create a k8s resource to manage your automation controller instance.

10.3. CONNECTING RESOURCE OPERATOR TO PLATFORM GATEWAY

To connect Resource Operator with platform gateway you need to create a k8s secret with the connection information for your automation controller instance.



NOTE

You can only create OAuth 2 Tokens for your own user through the API or UI, which means you can only configure or view tokens from your own user profile.

Procedure

To create an OAuth2 token for your user in the platform gateway UI:

1. Log in to Red Hat OpenShift Container Platform.
2. In the navigation panel, select **Access Management → Users**.

3. Select the username you want to create a token for.
4. Select **Tokens → Automation Execution**
5. Click **Create Token**.
6. You can leave **Applications** empty. Add a description and select **Read** or **Write** for the **Scope**.



NOTE

Make sure you provide a valid user when creating tokens. Otherwise, you will get an error message that you tried to issue the command without specifying a user, or supplying a username that does not exist.

10.4. CREATING A AUTOMATION CONTROLLER CONNECTION SECRET FOR RESOURCE OPERATOR

To make your connection information available to the Resource Operator, create a k8s secret with the token and host value.

Procedure

1. The following is an example of the YAML for the connection secret. Save the following example to a file, for example, **automation-controller-connection-secret.yml**.

```
apiVersion: v1
kind: Secret
metadata:
  name: controller-access
  type: Opaque
stringData:
  token: <generated-token>
  host: https://my-controller-host.example.com/
```

2. Edit the file with your host and token value.
3. Apply it to your cluster by running the **kubectl create** command:

```
kubectl create -f controller-connection-secret.yml
```

10.5. CREATING AN ANSIBLE JOB

Launch an automation job on automation controller by creating an AnsibleJob resource.

Procedure

1. Specify the connection secret and job template you want to launch.

```
apiVersion: tower.ansible.com/v1alpha1
kind: AnsibleJob
metadata:
  generateName: demo-job-1 # generate a unique suffix per 'kubectl create'
```

```
spec:
  connection_secret: controller-access
  job_template_name: Demo Job Template
```

2. Configure features such as, inventory, extra variables, and time to live for the job.

```
spec:
  connection_secret: controller-access
  job_template_name: Demo Job Template
  inventory: Demo Inventory # Inventory prompt on launch needs to be enabled
  runner_image: quay.io/ansible/controller-resource-runner
  runner_version: latest
  job_ttl: 100
  extra_vars: # Extra variables prompt on launch needs to be enabled
    test_var: test
  job_tags: "provision,install,configuration" # Specify tags to run
  skip_tags: "configuration,restart" # Skip tasks with a given tag
```



NOTE

You must enable prompt on launch for inventories and extra variables if you are configuring those. To enable **Prompt on launch**, within the automation controller UI: From the **Resources** → **Templates** page, select your template and select the **Prompt on launch** checkbox next to **Inventory** and **Variables** sections.

3. Launch a workflow job template with an AnsibleJob object by specifying the **workflow_template_name** instead of **job_template_name**:

```
apiVersion: tower.ansible.com/v1alpha1
kind: AnsibleJob
metadata:
  generateName: demo-job-1 # generate a unique suffix per 'kubectl create'
spec:
  connection_secret: controller-access
  workflow_template_name: Demo Workflow Template
```

10.6. CREATING A JOBTEMPLATE

- Create a job template on automation controller by creating a JobTemplate resource:

```
apiVersion: tower.ansible.com/v1alpha1
kind: JobTemplate
metadata:
  name: jobtemplate-4
spec:
  connection_secret: controller-access
  job_template_name: ExampleJobTemplate4
  job_template_project: Demo Project
  job_template_playbook: hello_world.yml
  job_template_inventory: Demo Inventory
```

